


Cross-dataset transfer — Model C → CCI v2 across two robust metrics



CCI v2 (DeBERTa-v3-large) beats its same-architecture baseline on HateXplain-Jewish on BOTH metrics (starred). On ToxiGen-Jewish the two methods are within seed noise — Model C edges AUROC/Macro-F₁ on the large backbones. Dots = seed-mean (5–6 seeds); AUROC is the 15-bin Mann-Whitney approximation.

